

LOOM VIDEO WALKTHROUGH GUIDE & SCRIPT

Scaler AI Engineer Screening | Candidate: Anurag Sajwan

1. OVERVIEW & PRE-RECORDING CHECKLIST

This guide provides a step-by-step recording plan and script for your Loom walkthrough video (target: <= 3 minutes). The goal is to demonstrate technical depth, system responsiveness, and clear reasoning under pressure to the Scaler evaluators.

Preparation Steps:

- [] Browser Tabs: Open Vercel Chat UI, Vapi Call logs, Cal.com Dashboard, and GitHub Repo (README.md).
- [] Terminal: Open a terminal window showing the fastapi logs in the background if running locally.
- [] Audio Check: Ensure your microphone is clear and background noise is minimized.
- [] Pace: Speak calmly and deliberately. Keep your cursor steady; avoid frantic page scrolling.

2. LOOM SCENE-BY-SCENE WALKTHROUGH SCRIPT

Scene 1: Introduction & Live Chat Demo (0:00 - 0:45)

Visual: Live deployed chat window (<https://scaler-ai-persona-pi.vercel.app>).

"Hi, I'm Anurag Sajwan. Today, I'll walk you through my Autonomous AI Persona, designed for the Scaler AI screening. Here is the live Next.js chat interface. Let's ask it a question about my background to show it is fully grounded in my resume: [Type: 'What is your experience with AI Agents?']... As you see, the response streams back instantly, pulling facts directly from my resume corpus. Let's book a slot: [Type: 'Book an interview']... The backend dynamically checks my real calendar and recommends only 3 open slots. Let's type 'ram, ram@gmail.com' and choose 9:30 AM. The booking is confirmed instantly."

Scene 2: Core Architecture Overview (0:45 - 1:20)

Visual: GitHub README.md showing the Mermaid architecture diagram.

"Here is the system architecture. We have a decoupled Next.js web client and a FastAPI backend. For voice calls, we use Vapi integrated with Deepgram STT, ElevenLabs TTS, and a GPT-4o-mini orchestrator, delivering natural, sub-2s responses with custom interruption handling. For chat RAG, we use NVIDIA NIM's nv-embedqa-e5-v5 embeddings stored in Pinecone serverless, and meta/llama-3.1-8b-instruct for lightning-fast token generation."

LOOM VIDEO WALKTHROUGH GUIDE & SCRIPT

Scaler AI Engineer Screening | Candidate: Anurag Sajwan

Scene 3: Hard Engineering Problems Solved (1:20 - 2:20)

Visual: Code editor showing backend/main.py, scroll to the book_meeting endpoint.

"The hardest problem I solved was preventing the LLM from hallucinating booking details. Eager models like Llama 3.1 8B often try to call the booking tool immediately by fabricating a date, time, and guest details, without asking the user. To solve this, I designed strict, stateful backend validation filters in FastAPI. First, the backend uses regular expressions to confirm a valid email address was explicitly entered by the user. Second, it parses the conversation history to verify that availability slots were actually checked and presented to the user. If either check fails, the backend raises a descriptive validation error, which forces the LLM to self-correct, ask the user for details, and follow the correct booking pipeline. This ensures a clean, bulletproof scheduling UX."

Scene 4: Latency Optimization & Evaluation Results (2:20 - 3:00)

Visual: Open the generated eval/EVAL_REPORT.pdf file.

"Finally, let's look at the evaluation report. By moving from Llama 3.1 70B to the 8B-Instruct model and implementing global HTTP connection pooling for embedding queries, we reduced first-token latency by over 90%, from 42 seconds to an average of 3.8 seconds. The voice agent's average response latency is 650 milliseconds, and the system achieves a 0% hallucination rate on the golden Q&A suite. Thanks for watching! All code is pushed and live."

3. KEY TALKING POINTS & TIPS FOR UPLOADING

- Highlight Honesty: Emphasize that the RAG model is strictly grounded and will report when it doesn't know an answer.
- Focus on System Design: Explain that you chose stateful backend filters over soft prompts because it provides absolute reliability.
- Keep it concise: Watch your timer during recording. If you go over 3 minutes, pause, delete, and re-record a slide.
- Trim if needed: Use Loom's built-in trim editor to cut out any long loading pauses or throat clears.